

## Homework Week 3 - Jon Workman

```
In [263]: from datetime import date
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_curve
%matplotlib inline
```

```
In [264]: df = pd.read_csv('seattleWeather_1948-2017.csv').dropna()
```

### Clean data and add feature

```
In [265]: df.RAIN = df.RAIN.astype(int)
df['DATE'] = pd.to_datetime(df['DATE'], format='%Y-%m-%d')
df['DAY'] = df['DATE'].apply(date.isoweekday)
```


```
In [266]: df.head()
```

Out[266]:


	DATE	PRCP	TMAX	TMIN	RAIN	DAY
0	1948-01-01	0.47	51	42	1	4
1	1948-01-02	0.59	45	36	1	5
2	1948-01-03	0.42	45	35	1	6
3	1948-01-04	0.31	45	34	1	7
4	1948-01-05	0.17	45	32	1	1

### Visualisation

```
In [267]: g = sns.heatmap(df_train[['RAIN', 'PRCP', 'DAY', 'TMIN', 'TMAX']].corr(),annot=True, fmt = ".2f",
cmap = "coolwarm")
```



```
In [268]: g = sns.lmplot(x='TMAX', y='PRCP', data=df, fit_reg=False, hue='RAIN', size=6)
```



## Split dataset

```
In [269]: start1 = pd.datetime(1950, 1, 1)
end1 = pd.datetime(2009, 12, 31)
start2 = pd.datetime(2010, 1, 1)
end2 = pd.datetime(2017, 12, 31)

df_train = df[(df.DATE >= start1) & (df.DATE <= end1)]
df_validation = df[(df.DATE >= start2) & (df.DATE <= end2)]
```

```
In [270]: df_train.shape
```

```
Out[270]: (21912, 6)
```

```
In [271]: df_validation.shape
```

```
Out[271]: (2905, 6)
```

## Train and test

```
In [272]: X_train, X_test, y_train, y_test = train_test_split(df_train.drop(['DATE', 'DAY', 'RAIN'], axis=
1), df_train['RAIN'],
test_size=0.30, random_state=0)
```

```
In [273]: logmodel = LogisticRegression()
logmodel.fit(X_train, y_train)
```

```
Out[273]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)
```

```
In [274]: logpredictions = logmodel.predict(X_test)
```


```
In [275]: print(accuracy_score(y_test, logpredictions))
```

```
0.92576817767
```

```
In [276]: print(confusion_matrix(y_test,logpredictions))
sns.heatmap(confusion_matrix(y_test, logpredictions) / len(y_test), cmap='Blues', annot=True)
```

```
[[3769   21]
 [ 467 2317]]
```


Out[276]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1a23553cc0>



```
In [277]: fpr, tpr, thresholds = roc_curve(y_test, logpredictions)

fig, ax = plt.subplots(1, figsize=(10, 5))
plt.plot(fpr, tpr, color='darkorange', label='Model Performace')
plt.plot([0, 1], [0, 1], color='gray', label='Random Performace')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Rain ROC Curve')
plt.legend(loc="lower right")
```

Out[277]: <matplotlib.legend.Legend at 0x1a23465198>



```
In [278]: print(classification_report(y_test,logpredictions))
```

	precision	recall	f1-score	support
0	0.89	0.99	0.94	3790
1	0.99	0.83	0.90	2784
avg / total	0.93	0.93	0.92	6574

## Validation

```
In [279]: X_train2, X_test2, y_train2, y_test2 = train_test_split(df_validation.drop(['DATE', 'DAY', 'RAIN'],axis=1),df_validation['RAIN'],
test_size=0.30,random_state=0)
```

```
In [280]: logmodel2 = LogisticRegression()  
logmodel2.fit(X_train2,y_train2)
```


```
Out[280]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,  
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,  
    verbose=0, warm_start=False)
```

```
In [281]: logpredictions2 = logmodel2.predict(X_test2)
```

```
In [282]: print(accuracy_score(y_test2,logpredictions2))  
  
0.857798165138
```


```
In [283]: print(confusion_matrix(y_test2,logpredictions2))  
sns.heatmap(confusion_matrix(y_test2, logpredictions2) / len(y_test2), cmap='Blues', annot=True  
)  
  
[[426  35]  
 [ 89 322]]
```

```
Out[283]: <matplotlib.axes._subplots.AxesSubplot at 0x1a23477908>
```



```
In [284]: fpr, tpr, thresholds = roc_curve(y_test2, logpredictions2)  
  
fig, ax = plt.subplots(1, figsize=(10, 5))  
plt.plot(fpr, tpr, color='darkorange', label='Model Performace')  
plt.plot([0, 1], [0, 1], color='gray', label='Random Performace')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('Rain ROC Curve')  
plt.legend(loc="lower right")
```

```
Out[284]: <matplotlib.legend.Legend at 0x1a2163eda0>
```



```
In [285]: print(classification_report(y_test2,logpredictions2))
```

	precision	recall	f1-score	support
0	0.83	0.92	0.87	461
1	0.90	0.78	0.84	411
avg / total	0.86	0.86	0.86	872